

CS3339 Machine Learning 2025 Fall Project

1 Introduction

In this project, you are required to complete a classification task with high-dimension feature data. The data come from an anonymous image classification dataset, in which each image is classified into one of the predefined 100 categories. As our course focus on statistical learning instead of feature extraction from raw data, we will provide the pre-extracted features from the dataset instead of the original images.

There are totally 19573 images for training, and we have another 10000 images for testing. Each image is represented as a 512 dim vector. You may simply load these features and their corresponding labels as a csv file. You may split a validation dataset with a preferred ratio by yourself.

Your objective is to train statistical learning models with data we provided, and achieve as high test accuracy as you can. Detailed descriptions are listed below.

2 Requirements

In this project, you should:

- Try at least **three** different statistical learning methods, and at least **two** of them are NOT deep learning. We recommend methods learnt in this course.
- Use at least **two** model assessment and selection methods to choose the best model and enhance the robustness and generalization ability of your models.
- Submit a report, your codes, and the results on test dataset as instructed in section 3.

In completing this task, you may want to perform:

- exploratory data analysis, including:
 - * understand the feature space and the distribution of classes,
 - * analyze the sparsity of the data and its implications;

- dimensionality reduction, including:
 - * apply techniques like PCA or SVD.
 - * discuss its impact on the performance and interpretability of models;
- model development, including:
 - * employ various statistical learning models (e.g., Logistic Regression, Linear SVM),
 - * experiment with different regularization techniques to handle overfitting;
- model evaluation: perform cross-validation to assess the generalization ability of the models.

Please be aware that although test accuracy (with its ranking) forms part of your score, it is not the only one. The detailed analysis, extensive experiments, a clear and well-written report, manual implementation of the algorithms based on basic libraries (e.g., numpy), reasonable modifications to standard algorithms and other highlights all contributes to a good project with high score. You may not wish to be stuck in the minor improvements of performance and ranking.

3 Submissions

You should submit your report and code via Canvas, and submit your results on test dataset with Kaggle platform. The report and code should be packaged to a zipped file, named with your student ID and name, with structure like:

```
ZhangSan-023XXXXXXXX.zip
├── Report.pdf
└── Code
    ├── Readme file
    └── Your code files here..
```

3.1 Report

Your report should describe how you complete this project, and it may include:

- the data processing techniques, and the reasons you choose them;
- a brief introduction of used models;
- the used model assessment and model selection methods, and how do you deal with the sparsity and high dimension of the data;
- how do you evaluate of your models;
- the conclusion;
- any other things you want to report in this project.

The report should be submitted as **a PDF file**, and recommend to have **six** or more pages. We suggest to use the **L^AT_EX**templates provided by top conferences and journals.

3.2 Code

Your code will be used for checking reproducibility. Therefore, your code should contain a Readme file, describing how to reproduce the results you submitted in kaggle and mentioned in the report. It could be a txt file for simple cases and recommended to be a PDF file if its complex.

3.3 Test Results

<https://www.kaggle.com/t/95bbe89d4bed41ecb3985cb25b93394f>

You need to submit your test results in this kaggle competition. You should submit your results before the submission deadline. You are also required to include the screen shots on the accuracy of methods used in your report.

4 Deadlines

Kaggle & Canvas Submission Deadline: January 11, 2026.